

华东师范大学数据科学与工程学院实验报告

课程名称：分布式编程模型与系统

指导教师：徐辰

上机实践名称：Hadoop 1.x 部署

年级：2019

姓名：赵煜硕

上机实践日期：

上机实践成绩：

学号：10195501415

2022.3.3

根据 lab1 注意事项，拉取 lab1 镜像到 lab2 创建如下云主机实例

基本信息

* 云账号名称:

UCloud-公有云/公有云 **UCloud**

* 资源类型:

云主机

* 所属资源集:

2019-分布式

实例信息

+ 新增实例

数量	地域	序号	配置信息 ①	操作系统	网络	预估价格	操作
1台	上海二可用区B	1	通用型 N12核4G	Ubuntu 18.04 64位	流量计费 20MB	0.33 元	

预估总价: 0.33 元 (可用金额: 9,982.80 元)

链接失败



刷新过后成功(欠刷新)

资源管理 / 云主机 / 创建

8

赵煜硕

创建UCloud云主机

数量	地域	序号	配置信息 ①	操作系统	网络	预估价格	创建状态
1台	上海二可用区B	1	通用型 N 2核 4G	Ubuntu 18.04 64位	流量计费 20MB	0.33 元	<div><div>✓ 开始</div><div>✓ 配额检测</div><div>✓ 创建云主机</div><div>✓ 手动同步公有</div><div>✓ 绑定云主机到</div><div>✓ 绑定云硬盘到</div><div>✓ 绑定EIP到资源</div><div>✓ 结束</div></div>

关闭

1台

上海二可用区B

1

通用型 N|2核|4G

Ubuntu 18.04 64位

流量计费|20MB

0.33 元

操作

预估总价: 0.33 元 (可用金额: 9,982.80 元)

华东师范大学数据科学与工程学院实验报告

课程名称：分布式编程模型与系统

年级：2019

上机实践成绩：

指导教师：徐辰

姓名：赵煜硕

学号：10195501415

上机实践名称：Hadoop 2.x 部署

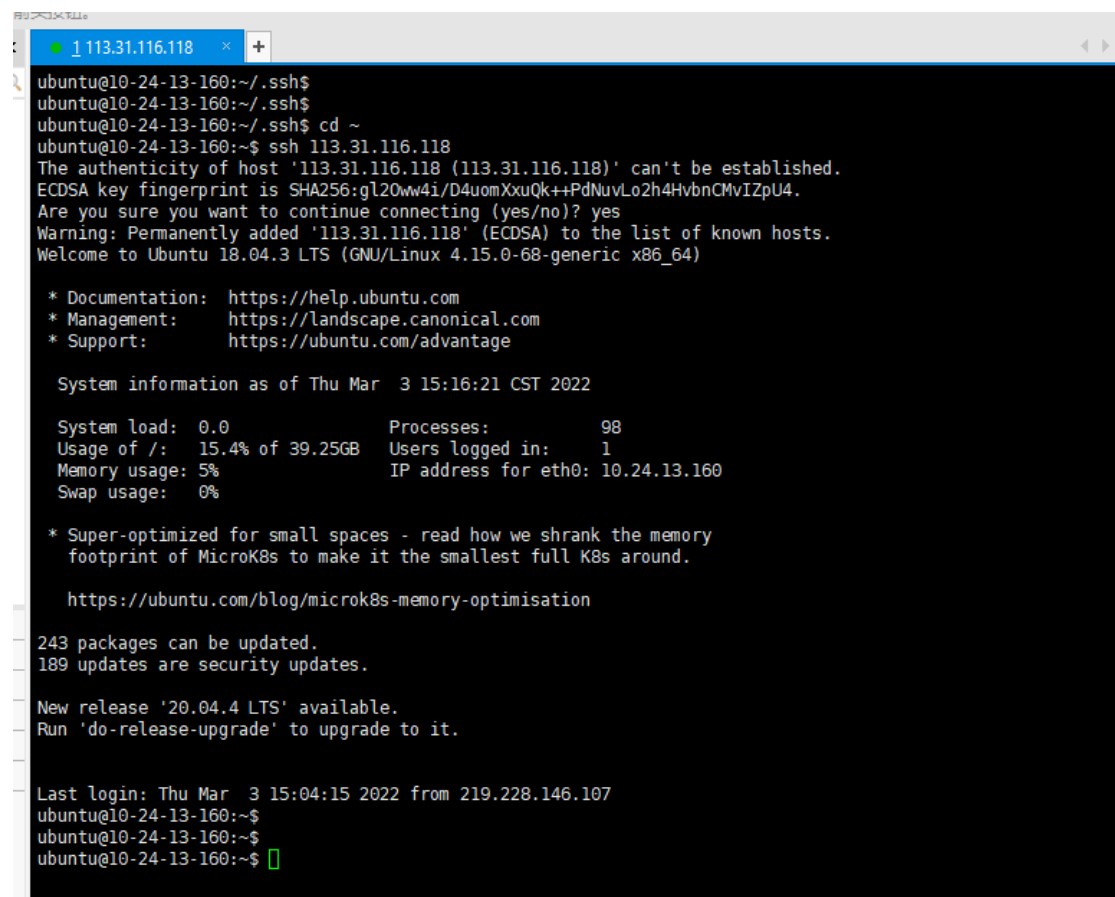
上机实践日期：

2022.3.10

Lab2

配置免密登陆失败

将 localhost 改为 127.0.0.1



```
1 113.31.116.118 x +
ubuntu@10-24-13-160:~/.ssh$
ubuntu@10-24-13-160:~/.ssh$
ubuntu@10-24-13-160:~/.ssh$ cd ~
ubuntu@10-24-13-160:~$ ssh 113.31.116.118
The authenticity of host '113.31.116.118 (113.31.116.118)' can't be established.
ECDSA key fingerprint is SHA256:gl20ww4i/D4uomXxuQk++PdNuvLo2h4HvbnCMvIZpU4.
Are you sure you want to continue connecting (yes/no)? yes
Warning: Permanently added '113.31.116.118' (ECDSA) to the list of known hosts.
Welcome to Ubuntu 18.04.3 LTS (GNU/Linux 4.15.0-68-generic x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:    https://landscape.canonical.com
 * Support:       https://ubuntu.com/advantage

System information as of Thu Mar  3 15:16:21 CST 2022

System load:  0.0               Processes:    98
Usage of /:   15.4% of 39.25GB   Users logged in: 1
Memory usage: 5%               IP address for eth0: 10.24.13.160
Swap usage:   0%

 * Super-optimized for small spaces - read how we shrank the memory
   footprint of MicroK8s to make it the smallest full K8s around.

https://ubuntu.com/blog/microk8s-memory-optimisation

243 packages can be updated.
189 updates are security updates.

New release '20.04.4 LTS' available.
Run 'do-release-upgrade' to upgrade to it.

Last login: Thu Mar  3 15:04:15 2022 from 219.228.146.107
ubuntu@10-24-13-160:~$
ubuntu@10-24-13-160:~$
ubuntu@10-24-13-160:~$
```

华东师范大学数据科学与工程学院实验报告

课程名称：分布式编程模型与系统

年级：2019

上机实践成绩：

指导教师：徐辰

姓名：赵煜硕

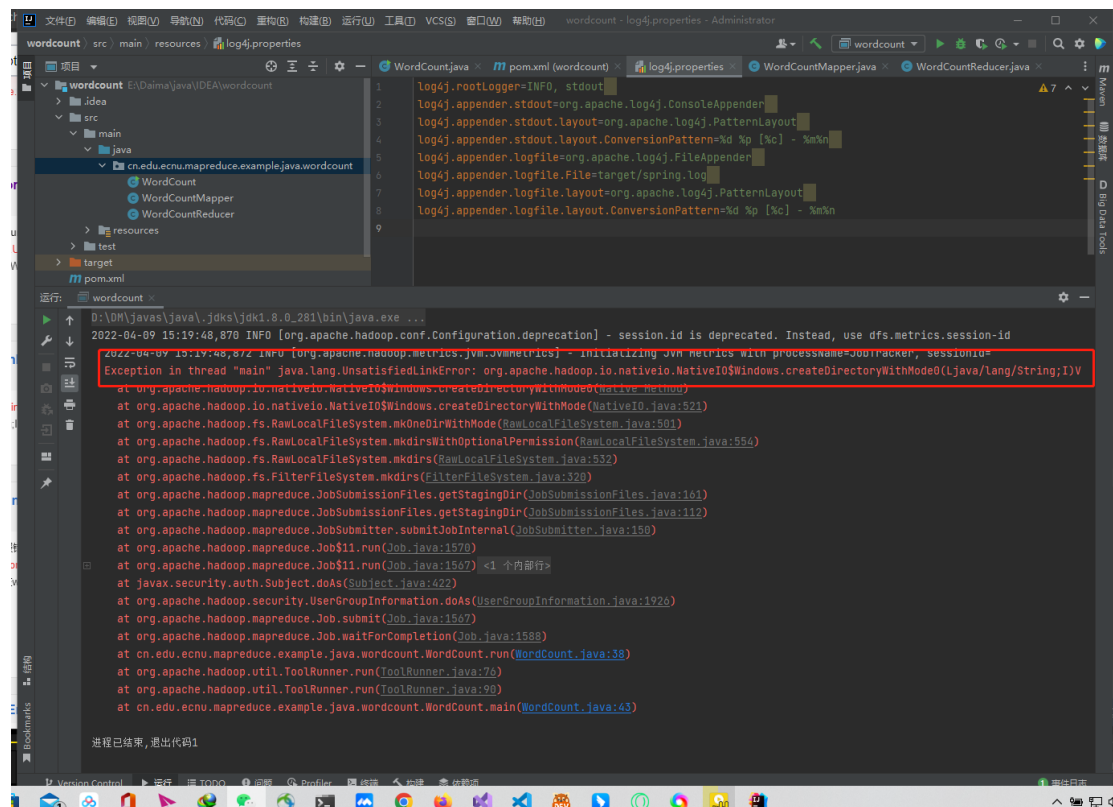
学号：10195501415

上机实践名称： MapReduce 2.x 编程

上机实践日期：

2022.3.24

Lab3



Windows 和 java 的 io 问题

Exception in thread "main" java.lang.UnsatisfiedLinkError: org.apache.hadoop.io.nativeio.NativeIO\$Windows.createDirectoryWithMode0(Ljava/lang/String;I)V

解决方法

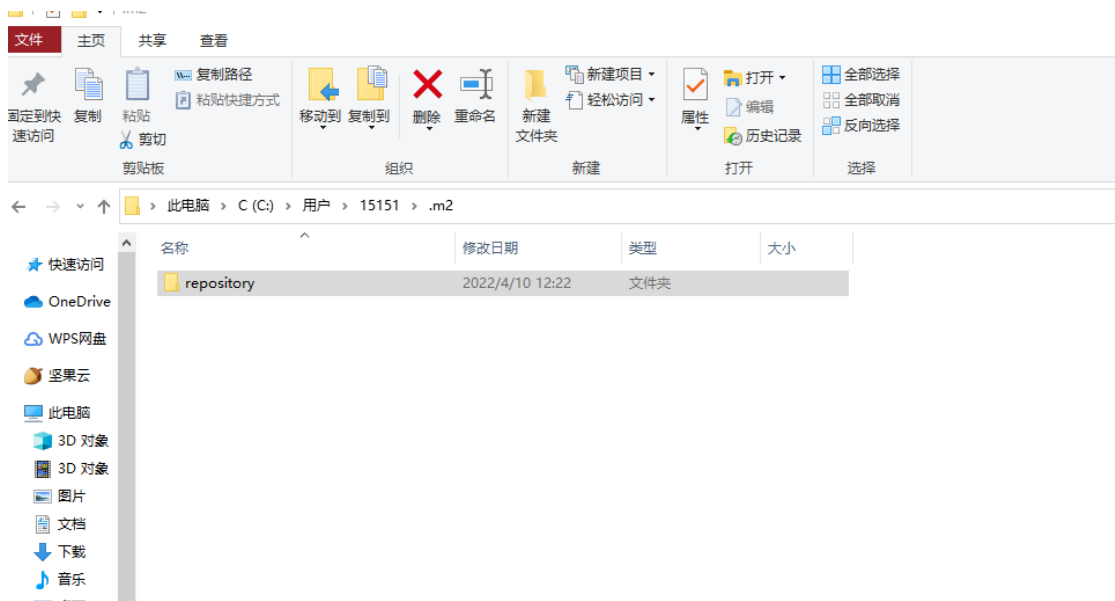
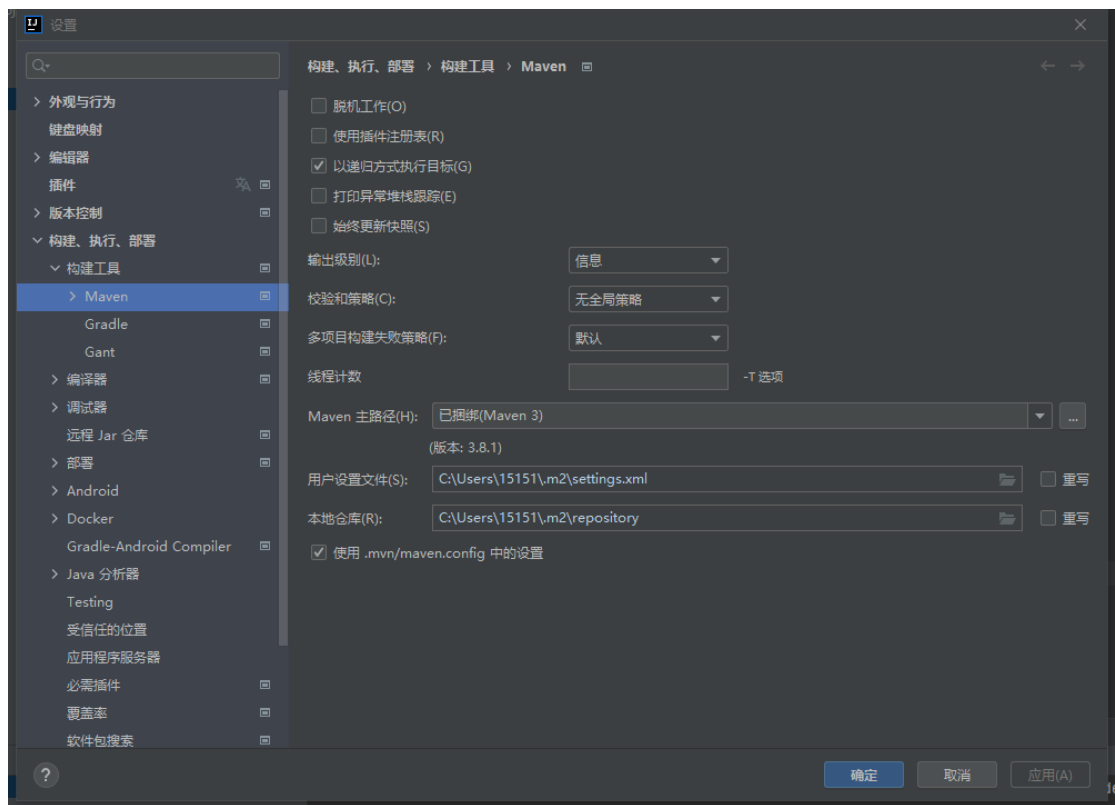
依赖版本改为 2.6.1 或者装 hadoop2.8.1 版本

IDEAbug

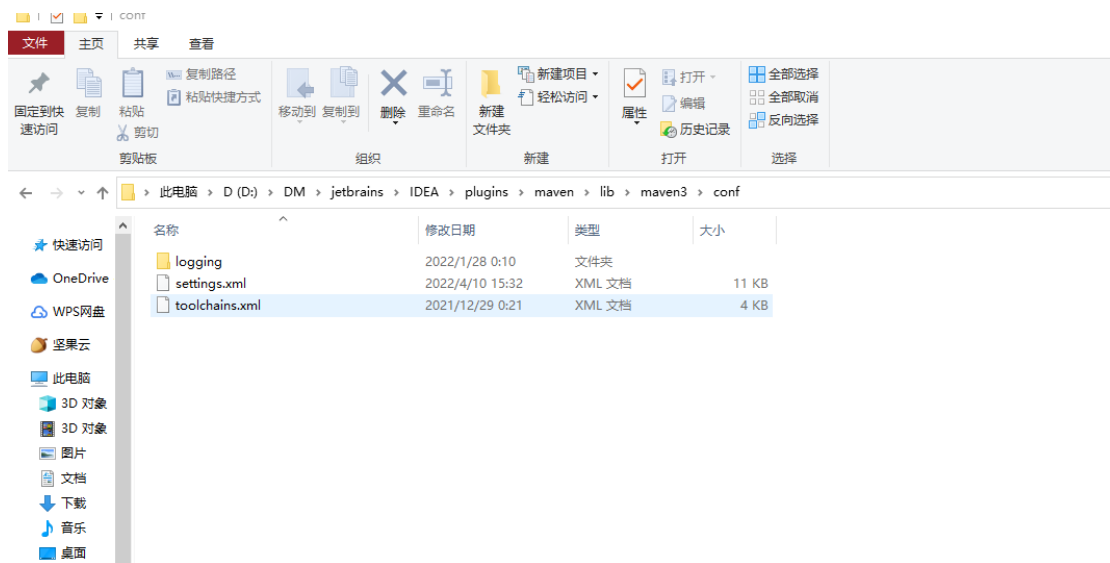
错误类型：

依赖一直在正在下载，毫无进度 <http://0.0.0.0>

Mavne 路径所指向的 setting.xml 不存在



默认使用以下图片中的路径



Mirror 问题

之前 mirror 指向 0.0.0.0,

是 idea 的问题, 问题来源于本机首次安装 IDEA 或者首次配置 maven。

在挂着 VPN 的状态下首次安装 IDEA, 会读取 ie 浏览器的配置会对 settings.xml 的 maven 的 mirror 进行重定向, 并且定位到 0.0.0.0

```

<!-- mirrors
| This is a list of mirrors to be used in downloading artifacts from remote repositories.
|
| It works like this: a POM may declare a repository to use in resolving certain artifacts.
| However, this repository may have problems with heavy traffic at times, so people have
| it to several places.
|
| That repository definition will have a unique id, so we can create a mirror reference for t
| repository, to be used as an alternate download site. The mirror site will be the preferre
| server for that repository.
|-->
<mirrors>
  <!-- mirror
  | Specifies a repository mirror site to use instead of a given repository. The repository th
  | this mirror serves has an ID that matches the mirrorOf element of this mirror. IDs are u
  | for inheritance and direct lookup purposes, and must be unique across the set of mirro
  |
  <mirror>
    <id>mirrorId</id>
    <mirrorOf>repositoryId</mirrorOf>
    <name>Human Readable Name for this Mirror.</name>
    <url>http://my.repository.com/repo/path</url>
  </mirror>
  -->
  <mirror>
    <id>alimaven</id>
    <name>aliyun maven</name>
    <url>http://maven.aliyun.com/nexus/content/groups/public/</url>
    <mirrorOf>central</mirrorOf>
  </mirror>
</mirrors>

<!-- profiles
| This is a list of profiles which can be activated in a variety of ways, and which can modify
| the build process. Profiles provided in the settings.xml are intended to provide local mac
| specific paths and repository locations which allow the build to work in the local environ
|
| For example, if you have an integration testing plugin - like cactus - that needs to know
| your Tomcat instance is installed, you can provide a variable here such that the variable
| dereferenced during the build process to configure the cactus plugin.
|
| As noted above, profiles can be activated in a variety of ways. One way - the activeProfil
| section of this document (settings.xml) - will be discussed later. Another way essentially
| relies on the detection of a system property, either matching a particular value for the p

```

```

<!-- mirrors
| This is a list of mirrors to be used in downloading artifacts from remote repositories.
|
| It works like this: a POM may declare a repository to use in resolving certain artifacts.
| However, this repository may have problems with heavy traffic at times, so people have
| it to several places.
|
| That repository definition will have a unique id, so we can create a mirror reference for t
| repository, to be used as an alternate download site. The mirror site will be the preferre
| server for that repository.
|-->
<mirrors>
  <!-- mirror
  | Specifies a repository mirror site to use instead of a given repository. The repository th
  | this mirror serves has an ID that matches the mirrorOf element of this mirror. IDs are u
  | for inheritance and direct lookup purposes, and must be unique across the set of mirro
  |
  <mirror>
    <id>mirrorId</id>
    <mirrorOf>repositoryId</mirrorOf>
    <name>Human Readable Name for this Mirror.</name>
    <url>http://my.repository.com/repo/path</url>
  </mirror>
  -->
  <mirror>
    <id>alimaven</id>
    <name>aliyun maven</name>
    <url>http://maven.aliyun.com/nexus/content/groups/public/</url>
    <mirrorOf>central</mirrorOf>
  </mirror>
</mirrors>

<!-- profiles
| This is a list of profiles which can be activated in a variety of ways, and which can modify
| the build process. Profiles provided in the settings.xml are intended to provide local mac
| specific paths and repository locations which allow the build to work in the local environ
|
| For example, if you have an integration testing plugin - like cactus - that needs to know
| your Tomcat instance is installed, you can provide a variable here such that the variable
| dereferenced during the build process to configure the cactus plugin.
|
| As noted above, profiles can be activated in a variety of ways. One way - the activeProfil
| section of this document (settings.xml) - will be discussed later. Another way essentially
| relies on the detection of a system property, either matching a particular value for the p

```